



FINAL-TERM PROJECT REPORT

Introduction To Data Science

Faculty

TOHEDUL ISLAM

Assistant Professor, Computer Science Dept.

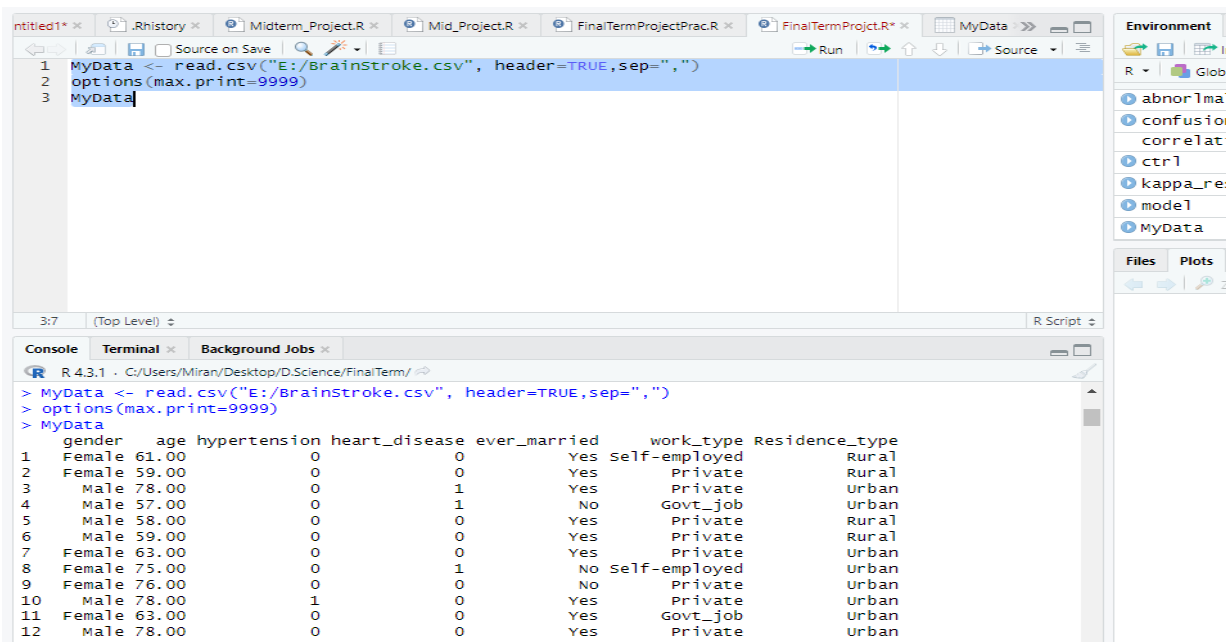
MD. ABDULLAH AL NASIR

19-41052-2

About Dataset: Here, the dataset named "Brain Stroke" has 202 instances and 11 attributes. Attributes are "Gender", "age", "hypertension", "heart_disease", "ever_married", "work_type", "Residence_type", "avg_glucose_level", "bmi", "smoking_status", "stroke".

Here, five attributes are categorical. Those are: "gender", "ever_married", "work_type", "Residence_type", "smoking_status". The attribute names: "age", "hypertension", "heart_disease", "avg_glucose_level", "bmi", "stroke" are numeric. The target attribute of this dataset is "stroke".

Step- 01: In the very first step, the dataset has been imported from the drive in a data frame named MyData using the read function. Using the options () function, we have imported all the data values with the help of max.print=9999.



The screenshot shows the R Studio interface. The top pane displays the R script with the following code:

```
1 MyData <- read.csv("E:/BrainStroke.csv", header=TRUE, sep=",")
2 options(max.print=9999)
3 MyData
```

The bottom pane shows the console output, which is a preview of the data frame:

	gender	age	hypertension	heart_disease	ever_married	work_type	Residence_type
1	Female	61.00	0	0	Yes	Self-employed	Rural
2	Female	59.00	0	0	Yes	Private	Rural
3	Male	78.00	0	1	Yes	Private	Urban
4	Male	57.00	0	1	No	Govt_job	Urban
5	Male	58.00	0	0	Yes	Private	Rural
6	Male	59.00	0	0	Yes	Private	Rural
7	Female	63.00	0	0	Yes	Private	Urban
8	Female	75.00	0	1	No	Self-employed	Urban
9	Female	76.00	0	0	No	Private	Urban
10	Male	78.00	1	0	Yes	Private	Urban
11	Female	63.00	0	0	Yes	Govt_job	Urban
12	Male	78.00	0	0	Yes	Private	Urban

Step- 02: The syntax of the function convertCategoricalToNumeric indicates the beginning of the code. This function expects MyData to be a data frame containing the dataset to be processed.

The function includes a list called categorical_cols that lists the column names that are required to be changed from categorical to numeric format. "Gender," "ever_married," "work_type," "Residence_type," and "smoking_status" are the columns listed here.

Each of the categorical columns listed in categorical_cols are iterated over in the code's loop. The actions listed below are carried out for each column:

- Factor Conversion:** The factor function converts the values of the column into factors. This step is essential because it assigns a numerical representation to each distinct value in the column.
- Numeric Conversion:** The as.numeric function is then used to convert the newly created factors into numeric values. This basically gives every individual factor value a numerical label.

```

1 MyData <- read.csv("E:/BrainStroke.csv", header=TRUE, sep=",")
2 options(max.print=9999)
3 MyData
4
5 convertCategoricalToNumeric <- function(MyData)
6 {
7   categorical_cols <- c("gender", "ever_married", "work_type", "Residence_type", "smoking_status")
8   for(col in categorical_cols)
9   {
10    MyData[[col]] <- as.numeric(factor(MyData[[col]]))
11  }
12  return(MyData)
13 }
14 MyData <- convertCategoricalToNumeric(MyData)
15 MyData

```

Console output:

```

R 4.3.1 - C:/Users/Miran/Desktop/D.Science/FinalTerm/
+ MyData[[col]] <- as.numeric(factor(MyData[[col]]))
+ }
+ return(MyData)
+ }
> MyData <- convertCategoricalToNumeric(MyData)
> MyData
  gender    age hypertension heart_disease ever_married work_type Residence_type
1      1  61.00           0              0           2         4             1
2      1  59.00           0              0           2         3             1
3      2  78.00           0              0           2         3             2
4      2  57.00           0              1           1         2             2
5      2  58.00           0              0           2         3             1
6      2  59.00           0              0           2         3             1
7      1  63.00           0              0           2         3             2

```

Step- 03: The function `colSums()` is used to calculate the column sums of a matrix or data frame. The function `is.na()` is used to identify null values (NA) in a vector, matrix, or data frame. To calculate the number of missing values in each column, the function used here is `colSums(is.na())`, keeping `MyData` as a parameter. Here we got that there are 0 null values in dataset.

```

1 MyData <- read.csv("E:/BrainStroke.csv", header=TRUE, sep=",")
2 options(max.print=9999)
3 MyData
4
5 convertCategoricalToNumeric <- function(MyData)
6 {
7   categorical_cols <- c("gender", "ever_married", "work_type", "Residence_type", "smoking_status")
8   for(col in categorical_cols)
9   {
10    MyData[[col]] <- as.numeric(factor(MyData[[col]]))
11  }
12  return(MyData)
13 }
14 MyData <- convertCategoricalToNumeric(MyData)
15 MyData
16
17 colSums(is.na(MyData))

```

Console output:

```

R 4.3.1 - C:/Users/Miran/Desktop/D.Science/FinalTerm/
197 70.15 29.75663 1 0
198 191.15 31.12417 3 0
199 95.02 31.79830 3 0
200 83.94 29.95130 3 0
201 83.75 29.09742 2 0
> colSums(is.na(MyData))
  gender    age hypertension heart_disease ever_married
0      0           0              0           0             0
  work_type Residence_type avg_glucose_level    bmi smoking_status
0      0           0              0           0             0
  stroke
0

```

Step- 04: Finding the outliers using Z-scores function. The `scale()` function is used to identify outliers by examining the standardized values of age, `abs()` returns the absolute value of a number. Outliers we found here are 1.32 3.00 6.00 7.00 0.48 5.00 1.88 1.08 1.80 3.00 1.32.

Now the outliers in the array named [age_outlier_indicates] are replaced with NA.

The screenshot shows an R script in RStudio. The script defines a function to convert categorical variables to numeric, then scales the 'age' variable. It identifies outliers where the absolute z-score is greater than 2. The outliers are then replaced with NA. The console output shows the z-scores and the resulting data frame with NA values for the outliers.

```

8 categorical_cols <- c(gender, ever_married, work_type, residence_type, smoking_status)
9 for(col in categorical_cols)
10 {
11   MyData[[col]] <- as.numeric(factor(MyData[[col]]))
12 }
13 return(MyData)
14 MyData <- convertCategoricalToNumeric(MyData)
15 MyData
16 colSums(is.na(MyData))
17
18 z_scores <- scale(MyData$age)
19 age_outlier_indices <- which(abs(z_scores) > 2)
20 age_outlier_values <- MyData$age[age_outlier_indices]
21 age_outlier_values
22
23 MyData$age[age_outlier_indices] <- NA
24 MyData
25

```

```

> z_scores <- scale(MyData$age)
> age_outlier_indices <- which(abs(z_scores) > 2)
> age_outlier_values <- MyData$age[age_outlier_indices]
> age_outlier_values
[1] 1.32 3.00 6.00 7.00 0.48 5.00 1.88 1.08 1.80 3.00 1.32
> MyData$age[age_outlier_indices] <- NA
> MyData
  gender age hypertension heart_disease ever_married work_type Residence_type avg_glucose_level
1     1   61             0              0           0         2             4             1             202.21
2     1   59             0              0           0         2             3             1             76.15
3     2   78             0              1           1         2             2             2             219.84
4     2   57             0              1           1         1             2             2             217.08

```

Step- 05: Another way to reduce outliers' values is to replace the most frequent or average value in the entire column. Using the table () function to calculate the most frequencies and which.max () to find the index of that value by using these functions together in the age attribute, we got the most frequent age, which is 78.

Now the data frame myFrequencyAgeData comparing with MyData and null values of age is replaced by 78.

The screenshot shows an R script in RStudio. The script identifies the most frequent age value (78) and replaces the outliers in the 'age' column with this value. It also creates a new data frame 'myFrequencyAgeData' which is a copy of 'MyData' with the outliers replaced. The console output shows the most frequent age value and the resulting data frame.

```

27
28 MyData$age[age_outlier_indices] <- NA
29 MyData
30
31 which(is.na(MyData$age))
32
33 most_frequent_age <- names(which.max(table(MyData$age)))
34 MyData$age[is.na(MyData$age)] <- most_frequent_age
35 most_frequent_age
36 myFrequencyAgeData = MyData
37 myFrequencyAgeData
38
39
40
41
42 MyData$age <- as.numeric(MyData$age)
43 MyData$stroke <- as.numeric(MyData$stroke)
44 MyData$stroke <- as.numeric(MyData$hypertension)
45
46

```

```

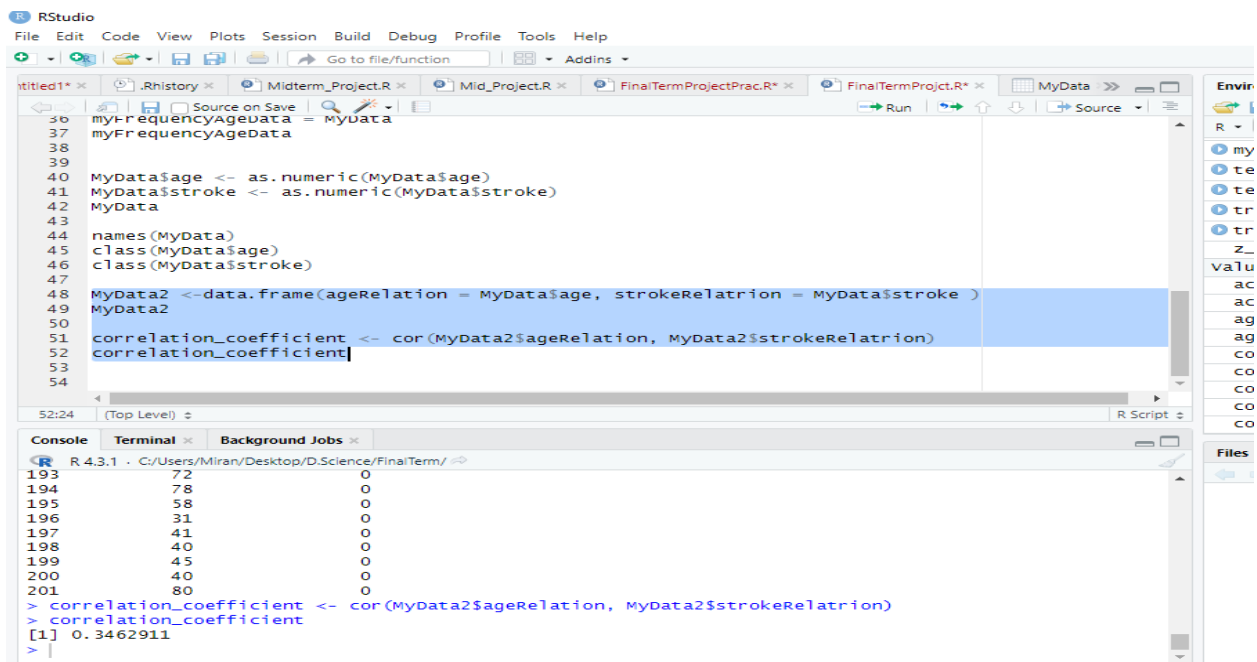
> most_frequent_age <- names(which.max(table(MyData$age)))
> MyData$age[is.na(MyData$age)] <- most_frequent_age
> most_frequent_age
[1] "78"
> myFrequencyAgeData = MyData
> myFrequencyAgeData
  gender age hypertension heart_disease ever_married work_type Residence_type avg_glucose_level
1     1   61             0              0           0         2             4             1             202.21
2     1   59             0              0           0         2             3             1             76.15
3     2   78             0              1           1         2             2             2             219.84
4     2   57             0              1           1         1             2             2             217.08

```

Step- 06: Two additional columns are being added to MyData2: The values from the age column of the original MyData data frame are entered into the first column, ageRelation. The values from the stroke column of the initial MyData data frame are placed in the second column, strokeRelation.

To find the correlation coefficient between two variables, use the `cor ()` function. In this instance, it will be applied to the MyData2 data frame's ageRelation and strokeRelation columns.

The variable `correlation_coefficient` contains the correlation calculation's outcome.



The screenshot shows the RStudio interface. The script editor contains the following code:

```
36 myFrequencyAgeData = MyData
37 myFrequencyAgeData
38
39
40 MyData$age <- as.numeric(MyData$age)
41 MyData$stroke <- as.numeric(MyData$stroke)
42 MyData
43
44 names(MyData)
45 class(MyData$age)
46 class(MyData$stroke)
47
48 MyData2 <- data.frame(ageRelation = MyData$age, strokeRelation = MyData$stroke)
49 MyData2
50
51 correlation_coefficient <- cor(MyData2$ageRelation, MyData2$strokeRelation)
52 correlation_coefficient
53
54
```

The console shows the output of the code:

```
R 4.3.1 > C:/Users/Miran/Desktop/D.Science/FinalTerm/
193      72      0
194      78      0
195      58      0
196      31      0
197      41      0
198      40      0
199      45      0
200      40      0
201      80      0
> correlation_coefficient <- cor(MyData2$ageRelation, MyData2$strokeRelation)
> correlation_coefficient
[1] 0.3462911
>
```

Step- 07: The R function `cor ()` is used to find the correlation matrix between variables in a dataset.

MyData: It is believed that this is a dataset (data frame or matrix) with a number of different variables. The `cor ()` function will determine how closely these variables are correlated.

"Pearson" is the correlation coefficient calculating method, according to this argument. The "Pearson" approach determines the Pearson correlation coefficient, which analyses the linear relationship between the variables. Another R function, **round ()**, rounds the values of a matrix or data frame to a predetermined number of digits.

`digits = 2:` This parameter indicates how many decimal places the correlation coefficients will be rounded to.

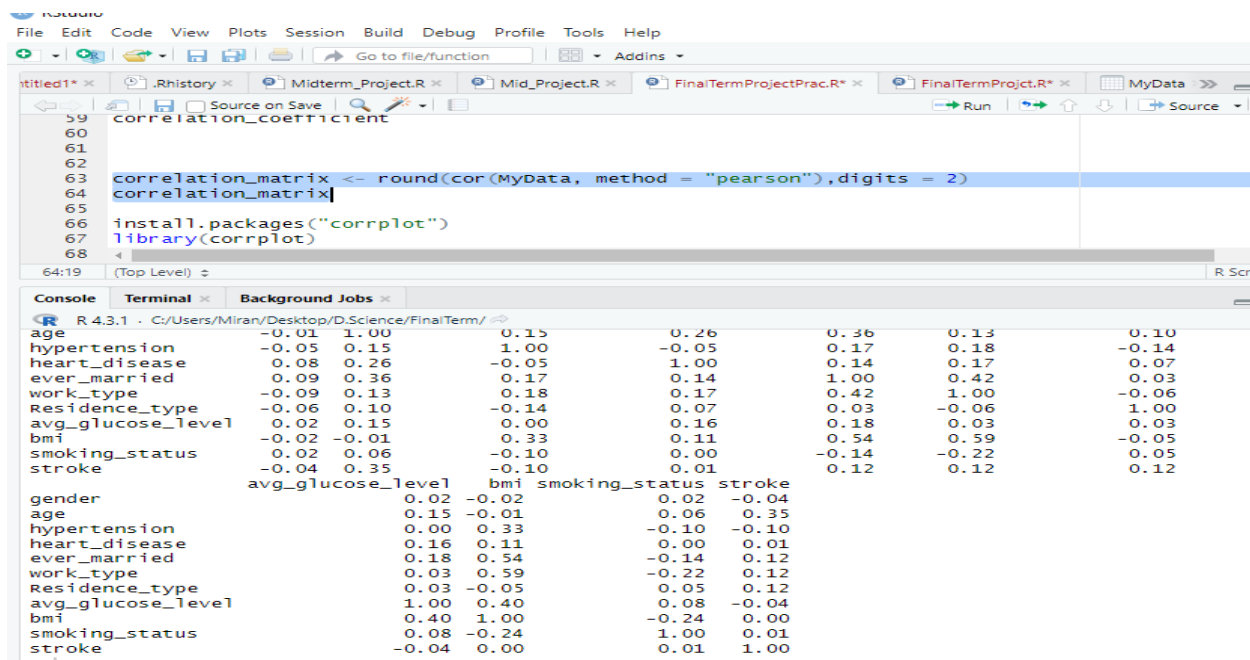
The correlation matrix will contain values between -1 and 1, where:

Values close to 1 indicate a strong positive correlation.

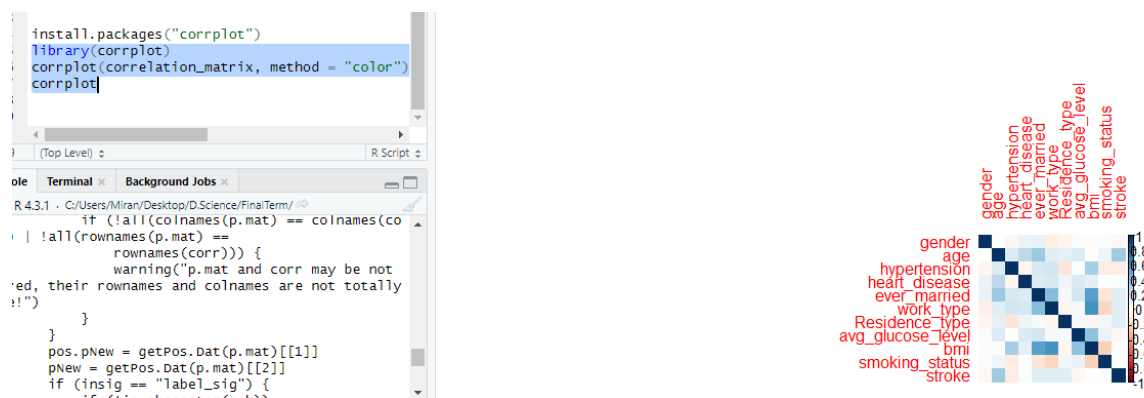
Values close to -1 indicate a strong negative correlation.

Values close to 0 indicate little or no linear correlation.

A strong positive correlation is found the output of this dataset.



Step- 08: The code `corrplot (correlation_matrix, method = "color")` generates a color-coded visualization of the correlation matrix stored in `correlation_matrix`. It uses the "corrplot" function, likely from the "corrplot" package, to create a plot where colors represent the strength and direction of correlations between variables.



Step- 09: `install. packages("class") library(class)` - The "class" package provides functions for k-nearest neighbors (KNN) classification and regression.

In this section we creating two data frame named training and test dataset. Here the code creates the `train_data` dataframe by extracting the "age" and "stroke" columns from the dataset, assuming there's one called `MyData`. The same columns are used to generate the test data dataframe as well.

The KNN algorithm's parameter `k`, which specifies how many neighbors to take into consideration when producing predictions, is set in this line. `K` in this case is set to 3.

The output ends with "0 0 0 0 0 0", indicating the predicted class labels for the last few data points in the test dataset.

Step- 10: `install.packages("caTools")` `library(caTools)`- "caTools" package, which provides various functions for data manipulation and splitting datasets.

The train_data dataframe has a total of 161 rows, according to this output. Based on the outcome of the splitting procedure, which allocated about 80% of the data to the training set, these rows were chosen from the original dataset.

The test_data dataframe has 40 rows in total, according to this output. Based on the outcome of the splitting procedure, which allocated about 20% of the data to the test set, these rows were also chosen from the original dataset.

```

58 install.packages("class")
59 library(class)
60
61 train_data <- MyData[c("age", "stroke")]
62 test_data <- data.frame(age = MyData$age, stroke = MyData$stroke)
63 k <- 3
64 knn_pred <- knn(train = train_data, test = test_data, cl = MyData$stroke, k = k)
65 knn_pred
66
67
68
69 install.packages("caTools")
70 library(caTools)
71
72 split <- sample.split(myFrequencyAgeData$stroke, SplitRatio = 0.8)
73
74
75 train_data <- myFrequencyAgeData[split, ]
76 nrow(train_data)
77 test_data <- myFrequencyAgeData[!split, ]
78 nrow(test_data)
79

```

```

R 4.3.1 - C:/Users/Miran/Desktop/D.Science/FinalTerm/
The downloaded binary packages are in
C:\Users\Miran\AppData\Local\Temp\Rtmp5W65Fo\downloaded_packages
> library(caTools)
> split <- sample.split(myFrequencyAgeData$stroke, SplitRatio = 0.8)
> train_data <- myFrequencyAgeData[split, ]
> nrow(train_data)
[1] 161
> test_data <- myFrequencyAgeData[!split, ]
> nrow(test_data)
[1] 40
>

```

Step – 11: This code performs k-fold cross-validation for a k-nearest neighbors (KNN) classification model on the 'MyData' dataset. It starts by randomly assigning instances to 'num_folds' (which is 10 in this case) different folds. Then, for each fold, it separates the data into training and testing sets. The KNN algorithm is trained on the training set's 'age' and 'stroke' columns, and then predictions are made for the 'stroke' values in the test set. The accuracy of predictions for each fold is computed by comparing KNN predictions to the actual 'stroke' values in the test set. The accuracy scores for all folds are stored in the 'accuracy_scores' array. This process helps assess the KNN model's performance across different data subsets and provides insight into its overall effectiveness.

```

83 num_folds <- 10
84
85 fold_indices <- sample(1:num_folds, size = nrow(MyData), replace = TRUE)
86 fold_indices
87
88
89
90 for (fold in 1:num_folds) {
91   train_indices <- which(fold_indices != fold)
92   test_indices <- which(fold_indices == fold)
93
94   train_data <- MyData[train_indices, c("age", "stroke")]
95   test_data <- MyData[test_indices, c("age", "stroke")]
96
97   knn_pred <- knn(train = train_data, test = test_data, cl = MyData$stroke[train_indices], k = k)
98
99   correct_predictions <- sum(knn_pred == MyData$stroke[test_indices])
100   total_predictions <- length(test_indices)
101   accuracy_scores[fold] <- correct_predictions / total_predictions
102 }
103
104 accuracy_scores

```

```

R 4.3.1 - C:/Users/Miran/Desktop/D.Science/FinalTerm/
correct_predictions <- sum(knn_pred == MyData$stroke[test_indices])
total_predictions <- length(test_indices)
accuracy_scores[fold] <- correct_predictions / total_predictions
}
accuracy_scores
[1] 0.9565217 0.9565217 1.0000000 0.9545455 1.0000000 0.8421053 0.9583333 0.9473684 0.8947368 0.9444444
library(caret)

```

Step-12: By the help of the R function table (), one can count the instances of various combinations of values in two or more categorical variables.

Actual = test_data\$stroke is specifying the values of the "Actual" column in the contingency table.

[illegible]

This balance between precision and recall depends on the specific problem and the trade-offs you are willing to make between minimizing false positives and false negatives.

```

88
89 precision <- conf_matrix[2, 2] / sum(conf_matrix[, 2])
90 precision
91 recall <- conf_matrix[2, 2] / sum(conf_matrix[2, ])
92 recall
93
93:1 (Top Level) ↕

```

Console Terminal × Background Jobs ×

```

R 4.3.1 - C:/Users/Miray/Desktop/D.Science/FinalTerm/
1] "Precision: 1"
1] "Recall: 0.85"
precision <- conf_matrix[2, 2] / sum(conf_matrix[, 2])
recall <- conf_matrix[2, 2] / sum(conf_matrix[2, ])
precision <- conf_matrix[2, 2] / sum(conf_matrix[, 2])
precision
1] 1
recall
1] 0.85

```

